

Monitoring, Analyzing, and Visualizing Distributed Computing Clusters

Peter Bui, John Stupka
Computer Science
{buij,stupkajr}@uwec.edu

1 Student Activities and Learning

As noted in the research narrative, the project will consist of the following components:

1. **Collector:** We will implement an application that periodically collects system and user information and sends it to an aggregator.
2. **Aggregator:** The aggregator will receive information from multiple collectors and combine them into a time-based database for efficient querying and replay.
3. **Visualization:** A web site will pull data from the database and visualize the results.
4. **Analytics:** A query language will be added to the database that will allow system administrators to analyze the large set of data for trends and patterns.

Our goal is to have a functional system by Spring 2015 and to use it on the BGSC and DPL clusters. If successful, we will try to publish at MICS (April 2015) and HTCondor Week (May 2015).

The student selected for this project is John Stupka, who is a Junior with experience in web development. He was chosen because of his deep interests in network and graph-based algorithms and his motivation to learn. In this project, John will learn more about distributed computing, databases, and visualization. He will gain practical skills in Python programming, web development, and system administration.

2 Faculty Mentoring

A description of my mentoring principles can be found in: “Using Clusters in Undergraduate Research: Distributed Animation Rendering, Photo Processing, and Image Transcoding.” The relevant principles are:

1. **Develop applications, not infrastructure.** Students are motivated by real-world applications that they can relate to and are interested in developing. Therefore, I work to ensure that the projects are realistic, but also interesting to the students.

2. **Utilize high-level frameworks.** Because software can become quite complex, I often encourage my students to use high-level frameworks that will allow them to become productive quickly so they can explore the problem domain. Sometimes this means using off-the-shelf software, while at other times this means providing the students with my own code or prototype.
3. **Recognize that nothing is straightforward.** All research projects encounter obstacles. As a mentor, I strive to be patient and recognize that while a problem may appear straightforward to me, it may be quite challenging to an undergraduate. My role is to provide the student with the appropriate background knowledge and guidance to not only find a solution, but learn how to find future solutions.
4. **Practice incremental development.** To support the process of research, I practice incremental development. This means I meet with students once a week to discuss their progress, any problems they are encountering, and how to develop solutions to these challenges. Additionally, my students record their research journey in a blog that I read before each meeting. Moreover, I have the students use a version control system, which allows me to track the progress of their code.

When it comes time for a paper or presentation, the students do the bulk of the writing. Afterwards, I apply the same principles of iterative development by having the students edit the publication based on my feedback until it has reached a satisfactory form. For talks and posters, I also have the students perform a live test run before presenting in public to ensure they effectively communicate their work.

3 History of Prior Activities

A list of all the papers and presentations produced from my collaborative student-faculty research at UW-Eau Claire is provided in the attached Faculty Vita. Highlights include the following:

1. One best paper award in the education, outreach, and training track at CLUSTER, a national high performance computing conference.
2. Two third-place student-written papers at the Midwest Instruction and Computing Symposium, a regional Computer Science conference.
3. One invited presentation at HTCondor Week, an annual workshop at UW-Madison.

Also listed are all of the students I have worked with in both collaborative research and as an independent study advisor.